

#Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

#Reading the dataset & storing it in a variable

```
health_df=pd.read_csv("health care diabetes.csv")
health_df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Project Task: Week 1

```
health_df.shape
```

```
(768, 9)
```

```
health_df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
768.000000				
mean	3.845052	120.894531	69.105469	20.536458
79.799479				
std	3.369578	31.972618	19.355807	15.952218
115.244002				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	1.000000	99.000000	62.000000	0.000000

```

0.000000
50%      3.000000  117.000000      72.000000      23.000000
30.500000
75%      6.000000  140.250000      80.000000      32.000000
127.250000
max      17.000000  199.000000      122.000000      99.000000
846.000000

```

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

Check for missing values & columns having no variance

```
health_df.isnull().any()
```

```

Pregnancies      False
Glucose           False
BloodPressure     False
SkinThickness     False
Insulin           False
BMI               False
DiabetesPedigreeFunction  False
Age               False
Outcome           False
dtype: bool

```

```
health_df.var()
```

```

Pregnancies      11.354056
Glucose           1022.248314
BloodPressure     374.647271
SkinThickness     254.473245
Insulin           13281.180078
BMI               62.159984
DiabetesPedigreeFunction  0.109779
Age               138.303046
Outcome           0.227483
dtype: float64

```

```
health_df['Glucose'].value_counts()
```

```

99      17
100     17
111     14
129     14

```

```

125      14
      ..
191       1
177       1
44        1
62        1
190       1
Name: Glucose, Length: 136, dtype: int64
health_df['BloodPressure'].value_counts()

```

```

70      57
74      52
78      45
68      45
72      44
64      43
80      40
76      39
60      37
0       35
62      34
66      30
82      30
88      25
84      23
90      22
86      21
58      21
50      13
56      12
52      11
54      11
75       8
92       8
65       7
85       6
94       6
48       5
96       4
44       4
100      3
106      3
98       3
110      3
55       2
108      2
104      2
46       2
30       2
122      1

```

```
95      1
102     1
61      1
24      1
38      1
40      1
114     1
```

```
Name: BloodPressure, dtype: int64
```

```
health_df['SkinThickness'].value_counts()
```

```
0      227
32     31
30     27
27     23
23     22
33     20
28     20
18     20
31     19
19     18
39     18
29     17
40     16
25     16
26     16
22     16
37     16
41     15
35     15
36     14
15     14
17     14
20     13
24     12
42     11
13     11
21     10
46      8
34      8
12      7
38      7
11      6
43      6
16      6
45      6
14      6
44      5
10      5
48      4
47      4
```

```

49      3
50      3
8       2
7       2
52      2
54      2
63      1
60      1
56      1
51      1
99      1
Name: SkinThickness, dtype: int64

```

```
health_df['Insulin'].value_counts()
```

```

0      374
105     11
130      9
140      9
120      8
...
73      1
171     1
255     1
52      1
112     1
Name: Insulin, Length: 186, dtype: int64

```

The dataset does not contain missing values(i.e. NaN) but presnce of 0 in features such as BP,Insulin,Glucose,etc. indicates the presence of Null values. So, here we will try to replace zero with np.NaN & later impute those null values.

```
health_df[['Glucose','BloodPressure','SkinThickness','Insulin','BMI']]
=health_df[['Glucose','BloodPressure','SkinThickness','Insulin','BMI']]
.replace(0,np.nan)
```

```
health_df.isnull().any()
```

```

Pregnancies      False
Glucose           True
BloodPressure     True
SkinThickness     True
Insulin           True
BMI               True
DiabetesPedigreeFunction  False
Age               False
Outcome           False
dtype: bool

```

Working with Missing Values & Imputation

```
health_df.sample(10)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
461	1	71.0	62.0	NaN	NaN	21.8
364	4	147.0	74.0	25.0	293.0	34.9
743	9	140.0	94.0	NaN	NaN	32.7
632	2	111.0	60.0	NaN	NaN	26.2
693	7	129.0	68.0	49.0	125.0	38.5
681	0	162.0	76.0	36.0	NaN	49.6
673	3	123.0	100.0	35.0	240.0	57.3
584	8	124.0	76.0	24.0	600.0	28.7
405	2	123.0	48.0	32.0	165.0	42.1
465	0	124.0	56.0	13.0	105.0	21.8

	DiabetesPedigreeFunction	Age	Outcome
461	0.416	26	0
364	0.385	30	0
743	0.734	45	1
632	0.343	23	0
693	0.439	43	1
681	0.364	26	1
673	0.880	22	0
584	0.687	52	1
405	0.520	26	0
465	0.452	21	0

```
health_df['Glucose'].median()
```

117.0

```
health_df['BloodPressure'].median()
```

72.0

```
health_df['SkinThickness'].median()
```

29.0

```
health_df['Insulin'].median()
```

125.0

```
health_df['BMI'].median()
```

32.3

```
health_df['Glucose']=health_df['Glucose'].fillna(117)
health_df['BloodPressure']=health_df['BloodPressure'].fillna(72)
health_df['SkinThickness']=health_df['SkinThickness'].fillna(29)
health_df['Insulin']=health_df['Insulin'].fillna(125)
health_df['BMI']=health_df['BMI'].fillna(32)
```

```
health_df.isnull().any()
```

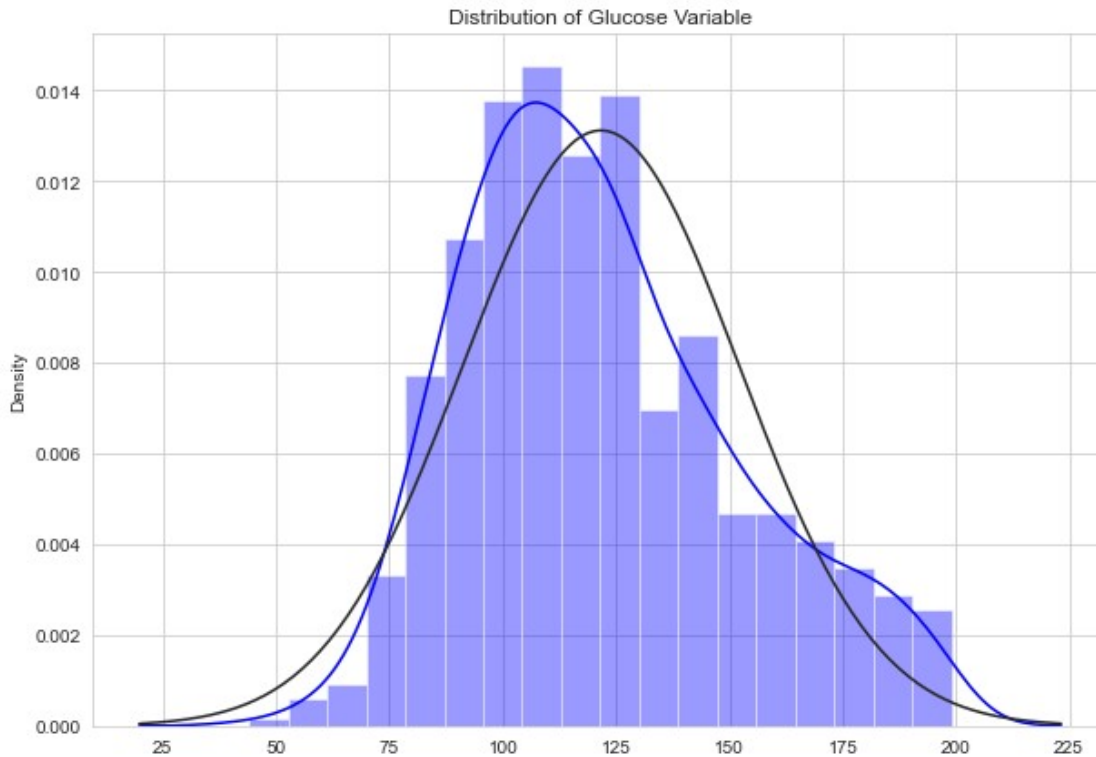
Pregnancies	False
Glucose	False
BloodPressure	False
SkinThickness	False
Insulin	False
BMI	False
DiabetesPedigreeFunction	False
Age	False
Outcome	False
dtype: bool	

Visualization to depict distribution of Variables

```
from scipy.stats import norm
plt.figure(figsize=(10,7))
sns.set_style('whitegrid')
```

```
sns.distplot(x=health_df['Glucose'],color='blue',fit=norm)
plt.title("Distribution of Glucose Variable")
plt.show()
```

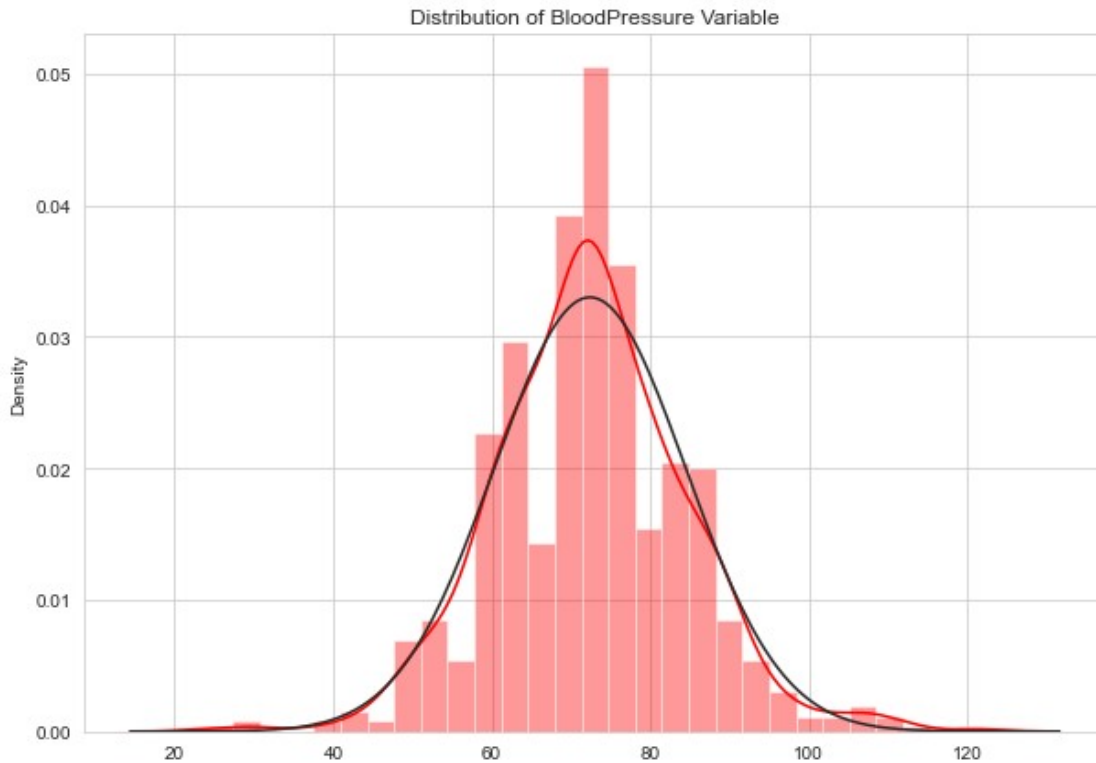
```
C:\Users\faiza\Anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```



```
from scipy.stats import norm
plt.figure(figsize=(10,7))
sns.set_style('whitegrid')
```

```
sns.distplot(x=health_df['BloodPressure'],color='red',fit=norm)
plt.title("Distribution of BloodPressure Variable")
plt.show()
```

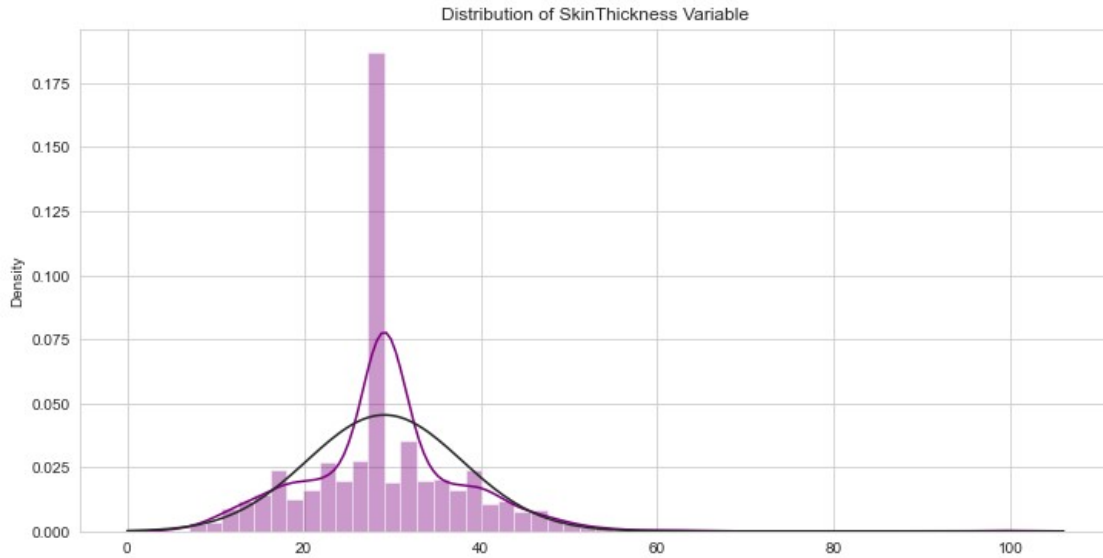
```
C:\Users\faiza\Anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

```
from scipy.stats import norm
plt.figure(figsize=(12,6))
sns.set_style('whitegrid')

sns.distplot(x=health_df['SkinThickness'],color='purple',fit=norm)
plt.title("Distribution of SkinThickness Variable")
plt.show()
```

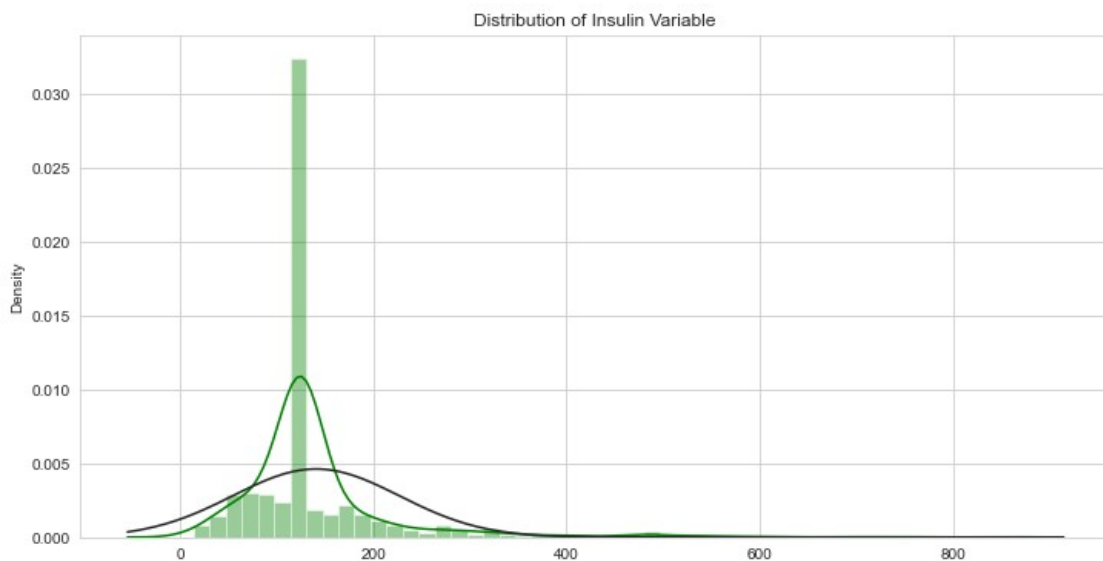
C:\Users\faiza\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



```
from scipy.stats import norm
plt.figure(figsize=(12,6))
sns.set_style('whitegrid')
```

```
sns.distplot(x=health_df['Insulin'],color='green',fit=norm)
plt.title("Distribution of Insulin Variable")
plt.show()
```

C:\Users\faiza\Anaconda3\lib\site-packages\seaborn\ distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



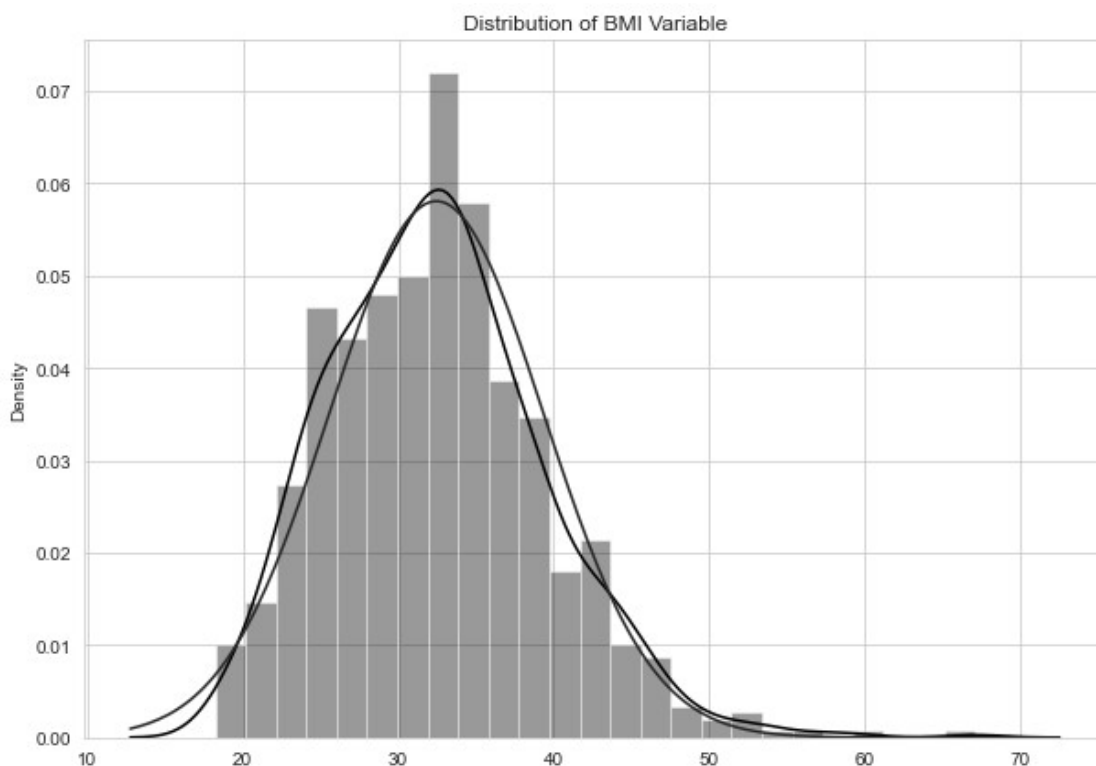
```

from scipy.stats import norm
plt.figure(figsize=(10,7))
sns.set_style('whitegrid')

sns.distplot(x=health_df['BMI'],color='black',fit=norm)
plt.title("Distribution of BMI Variable")
plt.show()

```

C:\Users\faiza\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



Above distplots show that there is slight deviation from Normal Distribution & data contains some Outliers.

Visual Exploration of Variables Using Histogram for possible values of Outcome(0,1)

```

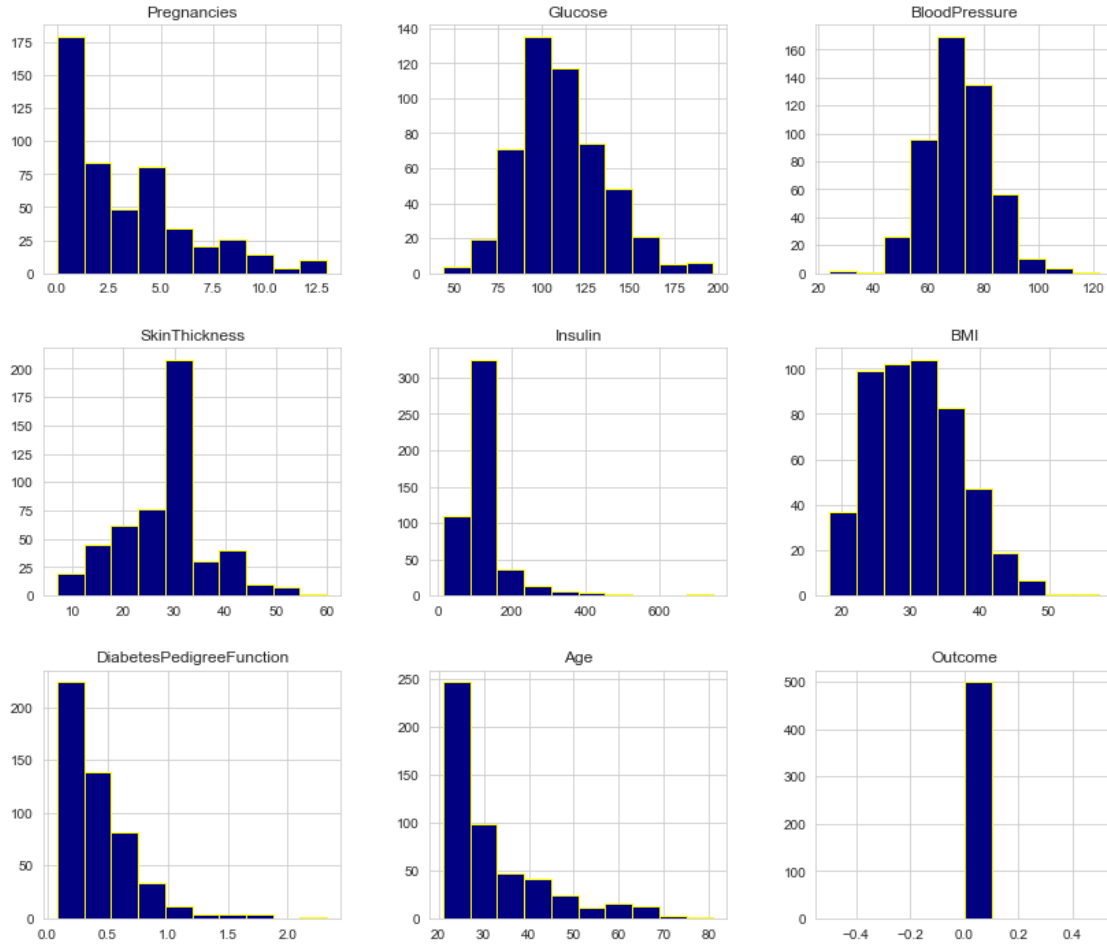
health_df.groupby('Outcome').hist(figsize=(14,12),color='navy',edgecolor='yellow')

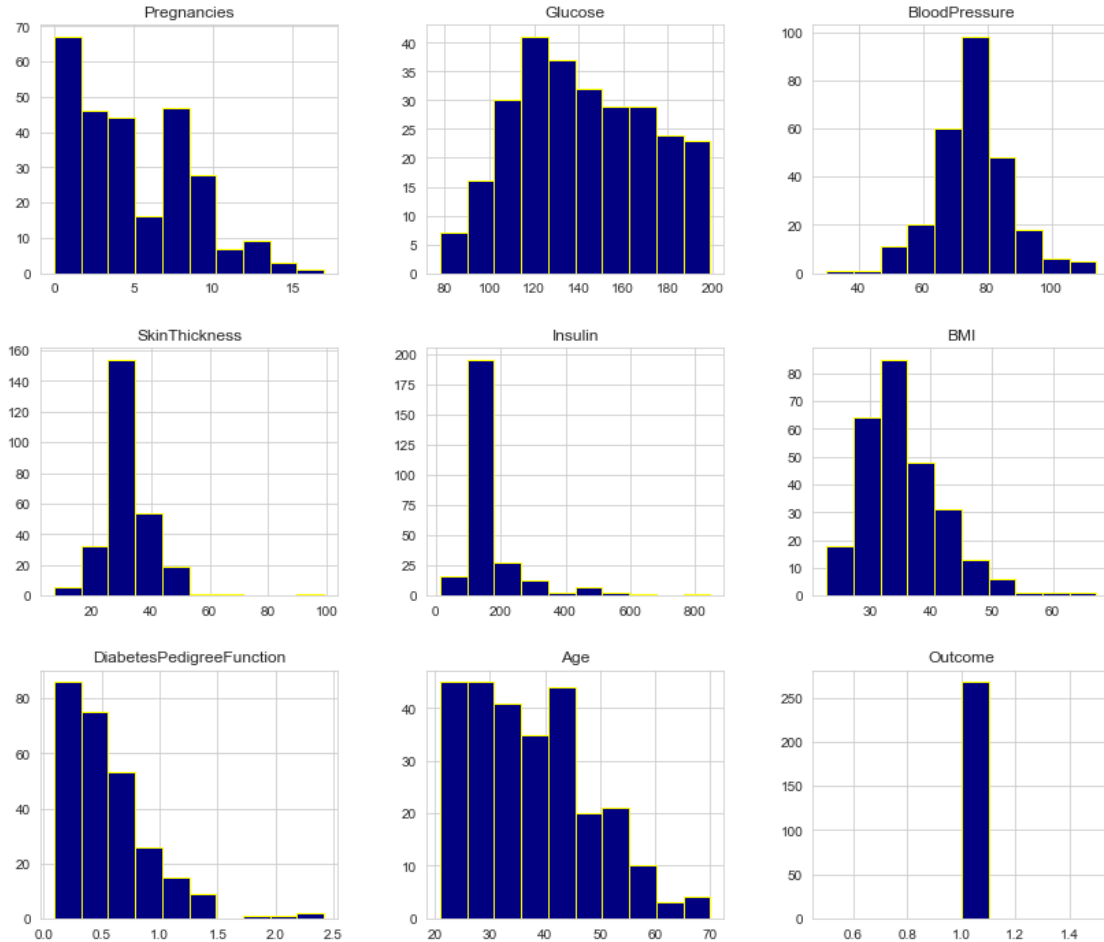
```

```

Outcome
0    [[AxesSubplot(0.125,0.670278;0.215278x0.209722...
1    [[AxesSubplot(0.125,0.670278;0.215278x0.209722...
dtype: object

```





Above code shows the frequency distribution of all the variables in the dataset for the possible cases of Non-Diabetic(0) & Diabetic(1)

Dtypes of Variables & Outcome

```
for col,val in health_df.iteritems():
    if(val.dtype=='int64'):
        if(val.nunique(>5):
            print(col," "+'Integer_Type-Continuous')

        else:
            print(col," "+'Integer_Type-Discrete or Categorical')

    elif (val.dtype=='float64'):
        if(val.nunique(>5):
            print(col," "+'Float_Type-Continuous')

        else:
            print(col," "+'Float_Type but Discrete or Categorical')

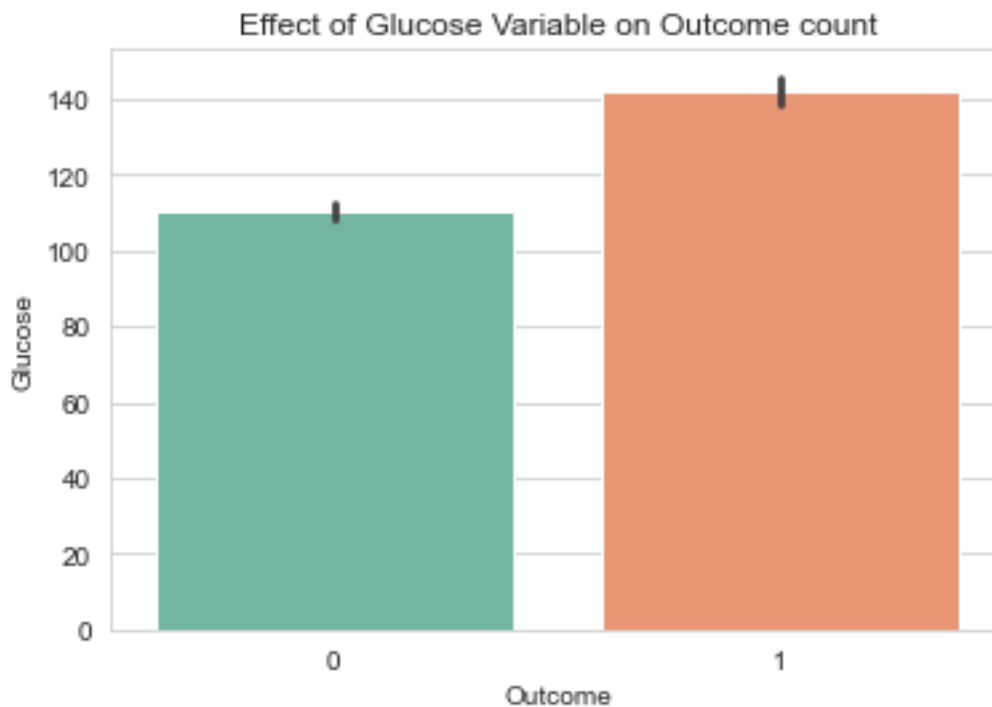
    else:
```

```
print(col, " "+'Object_DType-Categorical')
```

```
Pregnancies Integer_Type-Continuous  
Glucose Float_Type-Continuous  
BloodPressure Float_Type-Continuous  
SkinThickness Float_Type-Continuous  
Insulin Float_Type-Continuous  
BMI Float_Type-Continuous  
DiabetesPedigreeFunction Float_Type-Continuous  
Age Integer_Type-Continuous  
Outcome Integer_Type-Discrete or Categorical
```


Count of variables for outcome classes

```
sns.barplot(x='Outcome',y="Glucose",data=health_df,palette='Set2')  
plt.title("Effect of Glucose Variable on Outcome count")  
plt.show()
```




There seems to be a positive correlation between Glucose measure & Outcome.

```
sns.barplot(x='Outcome',y="BloodPressure",data=health_df,palette='Set1')  
plt.title("Effect of BloodPressure Variable on Outcome count")  
plt.show()
```



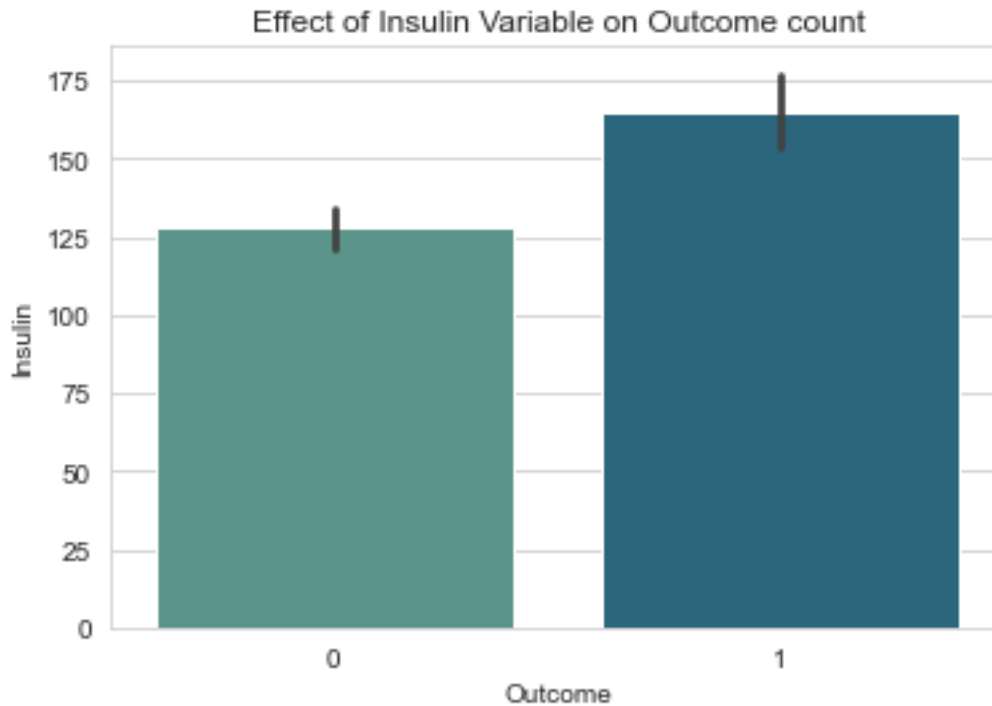
There seems to be a weak correlation between Glucose measure & Outcome.

```
sns.barplot(x='Outcome',y="SkinThickness",data=health_df,palette='magma')  
plt.title("Effect of skinThickness Variable on Outcome count")  
plt.show()
```




The correlation between skinThickness & Outcome doesn't seem to be significant.

```
sns.barplot(x='Outcome',y="Insulin",data=health_df,palette='crest')  
plt.title("Effect of Insulin Variable on Outcome count")  
plt.show()
```



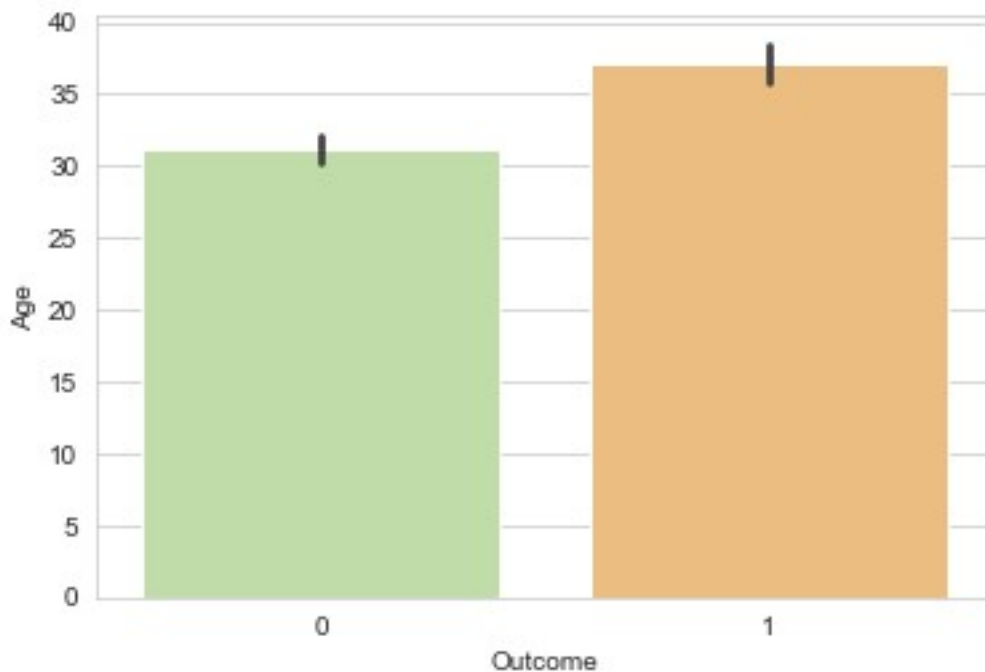
There correlation between Insulin & Outcome doesn't seem to be significant.

```
sns.barplot(x='Outcome',y="BMI",data=health_df,palette='Dark2')  
plt.title("Effect of BMI Variable on Outcome count")  
plt.show()
```

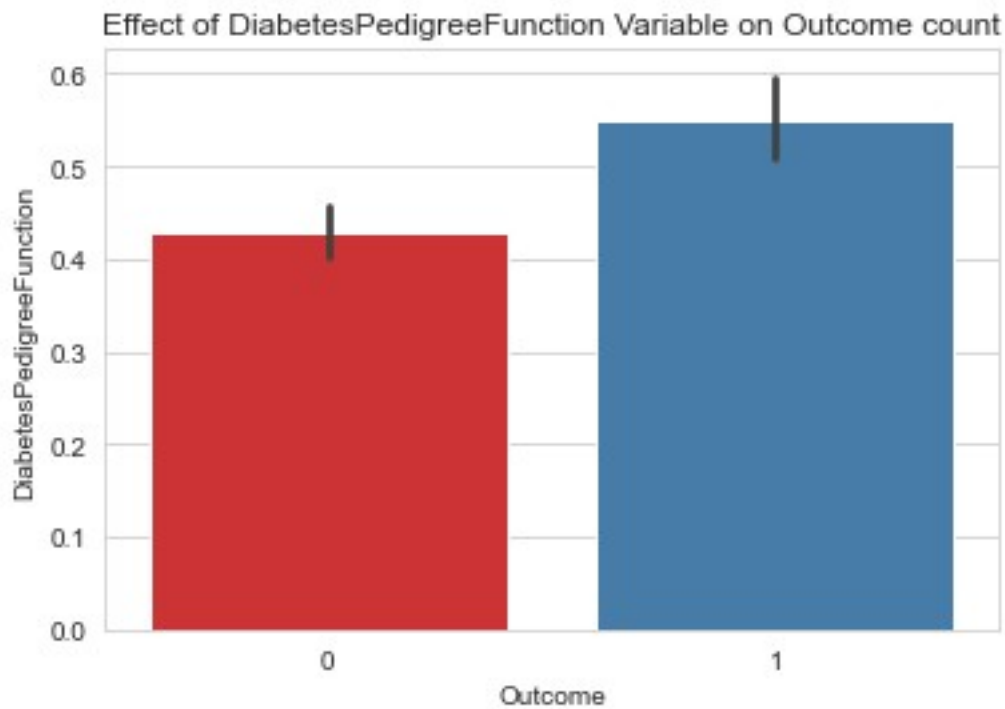
There seems to be a positive correlation between BMI measure & Outcome.

```
sns.barplot(x='Outcome',y="Age",data=health_df,palette='Spectral_r');
```



There seems to be a positive correlation between Age factor & Outcome.

```
sns.barplot(x='Outcome',y='DiabetesPedigreeFunction',data=health_df,palette='Set1');
plt.title("Effect of DiabetesPedigreeFunction Variable on Outcome count")
plt.show()
```




There seems to be a positive correlation between Glucose measure & Outcome.

Project Task: Week 2

Data Exploration

```
sns.countplot(x='Outcome',data=health_df,palette='PuRd');
```



```
health_df['Outcome'].value_counts()
```

```
0    500
```

```
1    268
```

```
Name: Outcome, dtype: int64
```

Here class is imbalanced ,where 65% data of Outcome Variable belongs to Class 1 & 35% data belongs to Class 0.

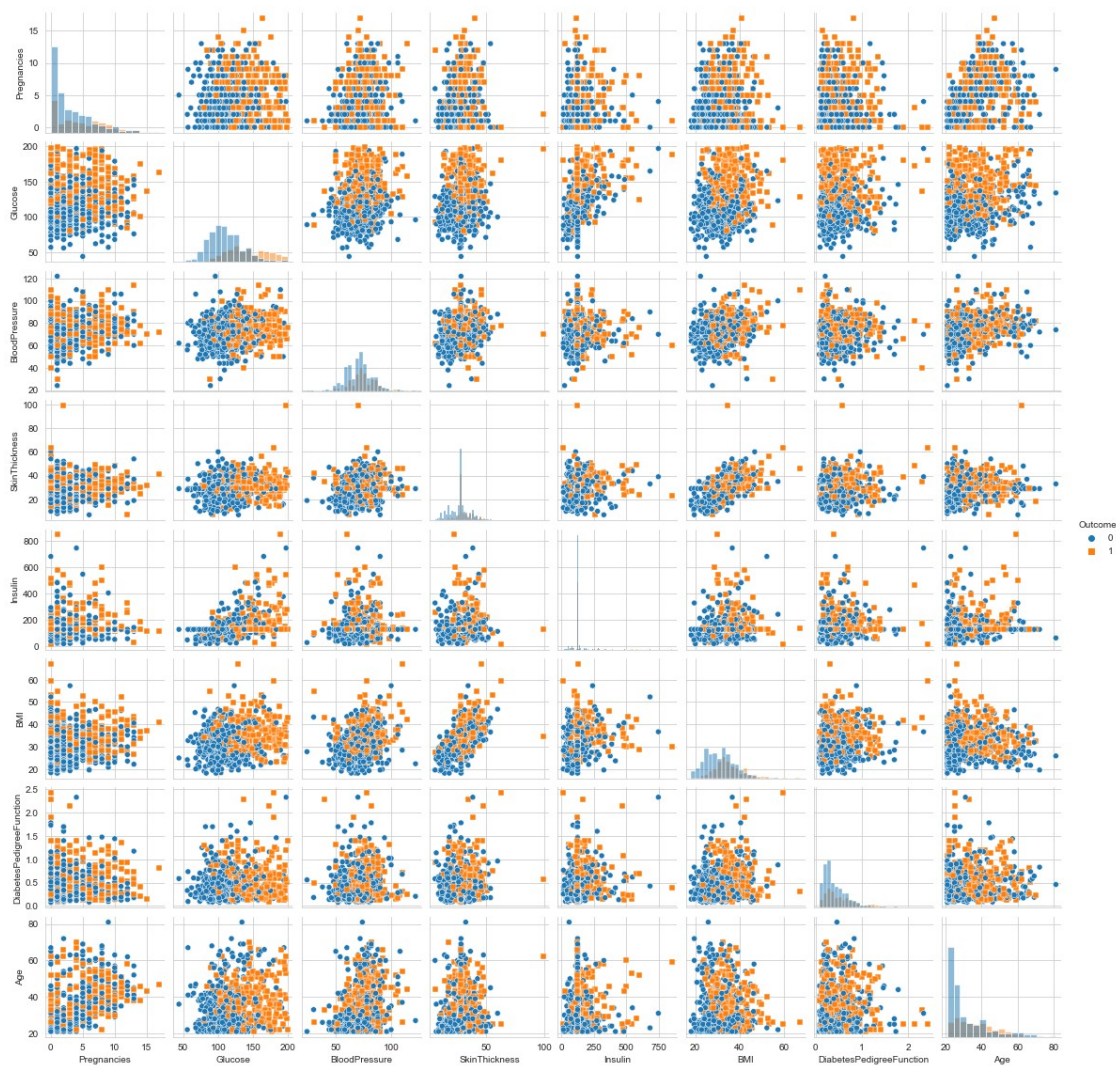
We will try to remove imbalance by Over-Sampling technique.

```
health_df.shape
```

```
(768, 9)
```

Relationship between different variables through Scatter plot employed using Pair plot

```
sns.pairplot(health_df,hue='Outcome',palette='tab10',height=2,diag_kin  
d='hist',markers=['o','s']);
```



out of all the plots SkinThickness & BMI seems to have some correlation.

Correlation Analysis

```
health_df.corr()
```

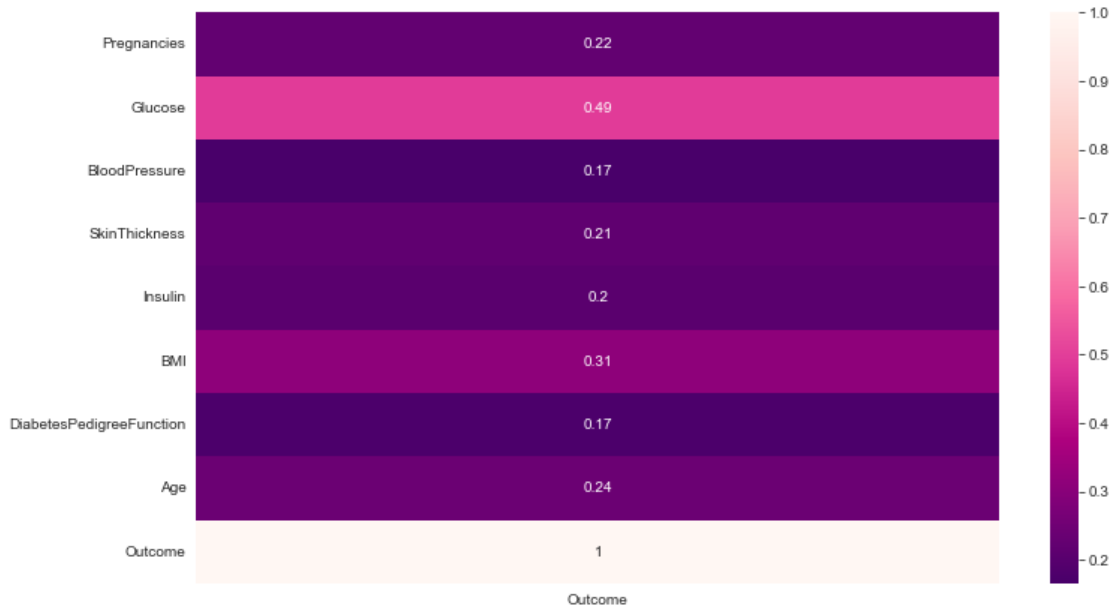
	Pregnancies	Glucose	BloodPressure
SkinThickness \			
Pregnancies	1.000000	0.128213	0.208615
0.081770			
Glucose	0.128213	1.000000	0.218937
0.192615			
BloodPressure	0.208615	0.218937	1.000000
0.191892			
SkinThickness	0.081770	0.192615	0.191892
1.000000			
Insulin	0.025047	0.419451	0.045363
0.155610			
BMI	0.021546	0.231400	0.281132

0.543275			
DiabetesPedigreeFunction	-0.033523	0.137327	-0.002378
0.102188			
Age	0.544341	0.266909	0.324915
0.126107			
Outcome	0.221898	0.492782	0.165723
0.214873			

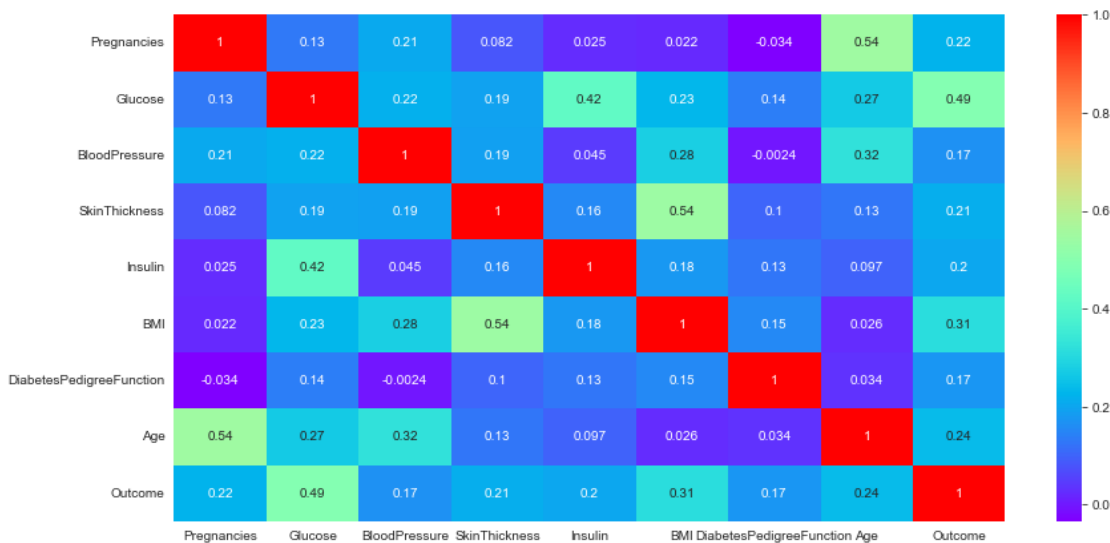
	Insulin	BMI	DiabetesPedigreeFunction
\			
Pregnancies	0.025047	0.021546	-0.033523
Glucose	0.419451	0.231400	0.137327
BloodPressure	0.045363	0.281132	-0.002378
SkinThickness	0.155610	0.543275	0.102188
Insulin	1.000000	0.180373	0.126503
BMI	0.180373	1.000000	0.153506
DiabetesPedigreeFunction	0.126503	0.153506	1.000000
Age	0.097101	0.025744	0.033561
Outcome	0.203790	0.312249	0.173844

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.266909	0.492782
BloodPressure	0.324915	0.165723
SkinThickness	0.126107	0.214873
Insulin	0.097101	0.203790
BMI	0.025744	0.312249
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
plt.figure(figsize=(12,7))
sns.heatmap(health_df.corr()[['Outcome']],annot=True, cmap='RdPu_r');
```



```
plt.figure(figsize=(14,7))
sns.heatmap(health_df.corr(),annot=True, cmap='rainbow');
```



Correlation analysis shows that there is no multi collinearity among the features.

It also seems that some independent variables may be insignificant for modelling purpose.

Outlier Detection & Removal

```
health_df.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
      'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

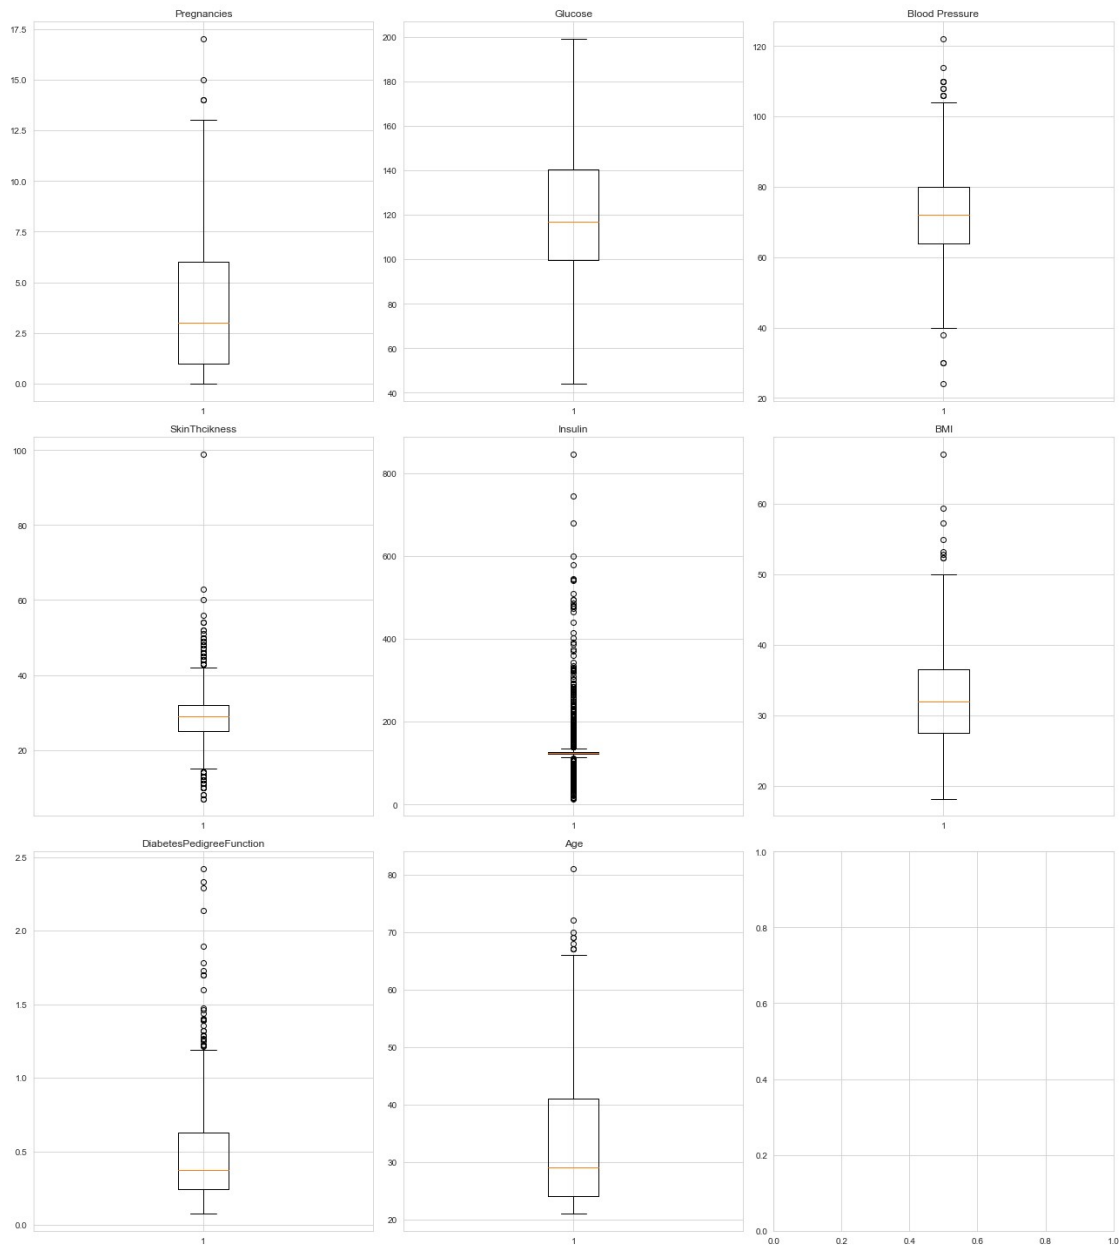
```

fig,ax=plt.subplots(3,3,figsize=(18,20))
ax[0,0].boxplot(health_df['Pregnancies'])
ax[0,1].boxplot(health_df['Glucose'])
ax[0,2].boxplot(health_df['BloodPressure'])
ax[1,0].boxplot(health_df['SkinThickness'])
ax[1,1].boxplot(health_df['Insulin'])
ax[1,2].boxplot(health_df['BMI'])
ax[2,0].boxplot(health_df['DiabetesPedigreeFunction'])
ax[2,1].boxplot(health_df['Age'])

ax[0,0].set_title('Pregnancies')
ax[0,1].set_title('Glucose')
ax[0,2].set_title('Blood Pressure')
ax[1,0].set_title('SkinThcikness')
ax[1,1].set_title('Insulin')
ax[1,2].set_title('BMI')
ax[2,0].set_title('DiabetesPedigreeFunction')
ax[2,1].set_title('Age')

plt.tight_layout()
plt.show();

```



```
health_df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	121.656250	72.386719	29.108073
std	3.369578	30.438286	12.096642	8.791221
min	0.000000	44.000000	24.000000	7.000000
25%	1.000000	99.750000	64.000000	25.000000


```

121.500000
50%      3.000000  117.000000      72.000000      29.000000
125.000000
75%      6.000000  140.250000      80.000000      32.000000
127.250000
max      17.000000  199.000000      122.000000      99.000000
846.000000

```

```

          BMI  DiabetesPedigreeFunction      Age      Outcome
count  768.000000      768.000000  768.000000  768.000000
mean    32.450911      0.471876   33.240885    0.348958
std      6.875366      0.331329   11.760232    0.476951
min     18.200000      0.078000   21.000000    0.000000
25%     27.500000      0.243750   24.000000    0.000000
50%     32.000000      0.372500   29.000000    0.000000
75%     36.600000      0.626250   41.000000    1.000000
max     67.100000      2.420000   81.000000    1.000000

```

#Creating a copy of Health_df

```

health=health_df.copy()
health.shape

```

(768, 9)

Project Task: Week 3

Check for significant features

```

import statsmodels.api as sm

features=health.drop(columns='Outcome')
target=health['Outcome']

sm_model=sm.OLS(target,features).fit()
sm_model
sm_model.summary()

<class 'statsmodels.iolib.summary.Summary'>
"""

```

OLS Regression Results

```

=====
=====
Dep. Variable:          Outcome      R-squared (uncentered):
0.502
Model:                  OLS         Adj. R-squared (uncentered):
0.496
Method:                 Least Squares   F-statistic:
95.66
Date:                   Sat, 08 Oct 2022   Prob (F-statistic):

```

```

1.28e-109
Time:                20:03:30   Log-Likelihood:
-417.96
No. Observations:    768   AIC:
851.9
Df Residuals:        760   BIC:
889.1
Df Model:             8

```

Covariance Type: nonrobust

[0.025 0.975]		coef	std err	t	P> t

Pregnancies		0.0247	0.005	4.613	0.000
0.014	0.035				
Glucose		0.0052	0.001	9.271	0.000
0.004	0.006				
BloodPressure		-0.0084	0.001	-7.227	0.000
-0.011	-0.006				
SkinThickness		-0.0009	0.002	-0.448	0.654
-0.005	0.003				
Insulin		-0.0001	0.000	-0.602	0.548
-0.001	0.000				
BMI		0.0072	0.003	2.724	0.007
0.002	0.012				
DiabetesPedigreeFunction		0.0925	0.047	1.984	0.048
0.001	0.184				
Age		0.0004	0.002	0.225	0.822
-0.003	0.004				
=====					
Omnibus:		195.899		Durbin-Watson:	
2.019					
Prob(Omnibus):		0.000		Jarque-Bera (JB):	
66.733					
Skew:		0.519		Prob(JB):	
3.23e-15					
Kurtosis:		1.995		Cond. No.	
675.					
=====					
=====					

Notes:

[1] R² is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is

```
correctly specified.  
"""
```

```
health.corr()['Outcome']*100
```

```
Pregnancies      22.189815  
Glucose           49.278240  
BloodPressure     16.572291  
SkinThickness     21.487322  
Insulin           20.379034  
BMI               31.224903  
DiabetesPedigreeFunction  17.384407  
Age               23.835598  
Outcome           100.000000  
Name: Outcome, dtype: float64
```

On analysing the p_value of features, we find that features like { Age, SkinThickness & Insulin } are statistically insignificant.

```
from sklearn.ensemble import ExtraTreesClassifier  
model=ExtraTreesClassifier()  
model.fit(features,target)
```

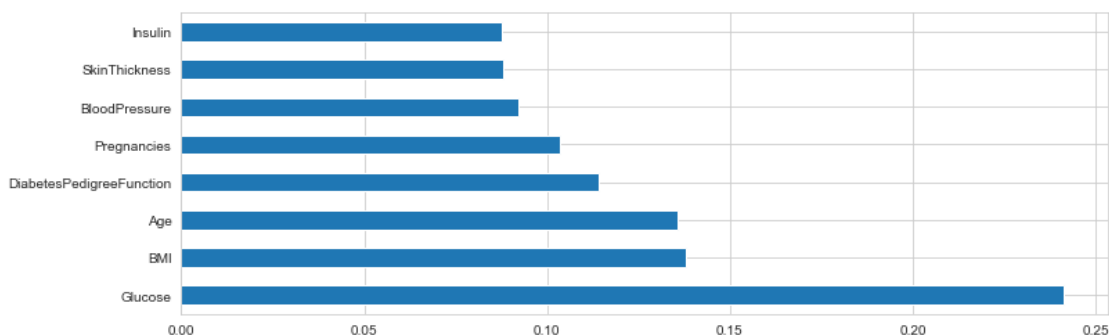
```
ExtraTreesClassifier()
```

```
print(model.feature_importances_)
```

```
[0.10344518 0.241295    0.09214672 0.08784809 0.08762691 0.13775837  
 0.11420148 0.13567825]
```

```
plt.figure(figsize=(12,4))  
ranked_features=pd.Series(model.feature_importances_,index=health.columns[:-1])  
ranked_features.nlargest(8).plot(kind='barh')
```

<AxesSubplot:>



```
health.corr()['Outcome'].sort_values(ascending=True)*100
```

```
BloodPressure      16.572291  
DiabetesPedigreeFunction  17.384407  
Insulin            20.379034
```

```

SkinThickness      21.487322
Pregnancies        22.189815
Age                23.835598
BMI                31.224903
Glucose            49.278240
Outcome            100.000000
Name: Outcome, dtype: float64

```

On analysing 1) the p_value of features using OLS Statsmodels we find that features like { Age, SkinThickness & Insulin} are statistically insignificant whereas 2) ExtraTreeClassifier predicts features like (BloodPressure, SkinThickness & Insulin) are least significant whereas 3) Correlation matrix describes BloodPressure, SkinThickness, Pregnancies, Insulin to be amongst least correlated to the outcome. So combining results from 1) & 2) and using correlation value,

We conclude that features like Insulin, SkinThickness are not significant.

Hence, we might not consider these features for model building.

```

features=features.drop(columns=['Insulin', 'SkinThickness'])
features.shape

(768, 6)

```

Outlier Detection & removal

```

q1=health_df['DiabetesPedigreeFunction'].quantile(0.25)
q3=health_df['DiabetesPedigreeFunction'].quantile(0.75)
iqr=q3-q1
max=1.5*iqr+ q3
min=q1-1.5*iqr
print(max,min)

1.2 -0.32999999999999996

np.where(health_df['DiabetesPedigreeFunction']>1.3)

(array([ 4, 12, 39, 45, 58, 147, 187, 228, 243, 259, 308, 330,
370,
       371, 395, 445, 593, 621, 622, 661], dtype=int64),)

np.where(health_df['DiabetesPedigreeFunction']<=-0.32)

(array([], dtype=int64),)

np.where(health_df['Pregnancies']>13)

(array([ 88, 159, 298, 455], dtype=int64),)

np.where(health_df['BMI']>60)

(array([177], dtype=int64),)

np.where(health_df['BloodPressure']>110)

```

```

(array([106, 691], dtype=int64),)
np.where(health_df['BloodPressure']<30)
(array([597], dtype=int64),)
outliers={4, 12, 39, 45, 58, 147, 187, 228, 243, 259, 308, 330,
370,
          371, 395, 593, 621, 622, 661, 88, 159, 298, 455}
outliers=list(outliers)
outliers

[259,
 4,
 455,
 330,
 395,
 12,
 593,
 147,
 661,
 88,
 159,
 228,
 39,
 298,
 45,
 621,
 622,
 370,
 243,
 308,
 371,
 58,
 187]

health=health.drop(outliers,axis=0)
health.shape

(745, 9)

```

Handling Class Imbalanced Data & Over_Sampling

```

health=health.drop(columns=['Insulin','SkinThickness'])
health.shape

(745, 7)

health['Outcome'].value_counts(normalize=True)*100

0    65.771812
1    34.228188
Name: Outcome, dtype: float64

```

```
health['Outcome'].value_counts()
```

```
0    490
```

```
1    255
```

```
Name: Outcome, dtype: int64
```

It is evident that Outcome has imbalanced data as 253 observations belong to Class 0 (i.e 34%)

while 489 observations belong to Class 1(i.e 66%)

Since the size of the Dataset is small, we can consider :

1)Over_Sampling Technique to balance the data

First we will divide the dataset into two parts, one for training & validation and other for testing. So, first we will use Train_test_split to divide the dataset & then use SKF for training & validation.

```
health.shape
```

```
(745, 7)
```

```
data=health.drop(columns='Outcome')
```

```
target=health['Outcome']
```

Splitting the dataset into 1)Train & Validation Set and 2)Test Set

```
from sklearn.model_selection import train_test_split
```

```
x, x_test, y, y_test = train_test_split(data, target, test_size = 0.2,  
random_state = 1)
```

```
print(x.shape, x_test.shape, y.shape, y_test.shape)
```

```
(596, 6) (149, 6) (596,) (149,)
```

```
y.value_counts()
```

```
0    391
```

```
1    205
```

```
Name: Outcome, dtype: int64
```

Oversampling by SMOTE

```
from imblearn.over_sampling import SMOTE
```

```
smote=SMOTE(sampling_strategy='minority')
```

```
x_os,y_os=smote.fit_resample(x,y)
```

```
y_os.value_counts()
```

```
0    391
```

```
1    391
```

```
Name: Outcome, dtype: int64
```

Both classes have the same number of count after Over sampling ,hence the data is balanced.

Stratified K-Fold CV

```
from sklearn.model_selection import StratifiedKFold
skf=StratifiedKFold(n_splits=5)
skf

StratifiedKFold(n_splits=5, random_state=None, shuffle=False)

from sklearn.neighbors import KNeighborsClassifier
model1=KNeighborsClassifier(n_neighbors=5, p=2, metric='minkowski')

from sklearn.linear_model import LogisticRegression
model2=LogisticRegression(solver = 'newton-cg', C = 1000,
penalty='l2')

from sklearn.ensemble import RandomForestClassifier
model3=RandomForestClassifier(n_estimators=50, max_depth=40)

from sklearn.svm import SVC
model4=SVC()

from sklearn.naive_bayes import GaussianNB
model5=GaussianNB(var_smoothing=1e-07)

from xgboost import XGBClassifier
model6=XGBClassifier(base_score=0.25, booster='gbtree',
                    colsample_bylevel=1, colsample_bynode=1,
                    colsample_bytrees=1, n_estimators=50)

from sklearn.model_selection import cross_val_score

def get_scores(model):
    score=cross_val_score(model, x_os, y_os, cv=skf)
    mean_score=np.mean(score)
    print('Model: ', model)
    return mean_score

get_scores(model1)
Model:  KNeighborsClassifier()
0.7763187979748489

get_scores(model2)
Model:  LogisticRegression(C=1000, solver='newton-cg')
0.7417034133594643

get_scores(model3)
Model:  RandomForestClassifier(max_depth=40, n_estimators=50)
0.826139147476727
```

```
get_scores(model4)
```

```
Model: SVC()
```

```
0.7276171811203659
```

```
get_scores(model5)
```

```
Model: GaussianNB(var_smoothing=1e-07)
```

```
0.7365997060264576
```

```
get_scores(model6)
```

```
Model: XGBClassifier(base_score=0.25, booster='gbtree',  
callbacks=None,  
                      colsample_bylevel=1, colsample_bynode=1,  
                      colsample_bytree=1,  
                      early_stopping_rounds=None, enable_categorical=False,  
                      eval_metric=None, gamma=None, gpu_id=None,  
                      grow_policy=None,  
                      importance_type=None, interaction_constraints=None,  
                      learning_rate=None, max_bin=None,  
                      max_cat_to_onehot=None,  
                      max_delta_step=None, max_depth=None, max_leaves=None,  
                      min_child_weight=None, missing=nan,  
                      monotone_constraints=None,  
                      n_estimators=50, n_jobs=None, num_parallel_tree=None,  
                      predictor=None, random_state=None, reg_alpha=None,  
                      reg_lambda=None, ...)
```

```
0.8094888126735261
```

Comparing the performance of KNN Algorithm with five other Algorithms using Stratified K-Fold CV score, KNN is 3rd ranked model in terms of accuracy while Random Forest gives the best result & Naive Bayes being the worst performer.

Project Task: Week 4

Model Building Using KNeighbours Classifier

```
from sklearn.neighbors import KNeighborsClassifier  
model1=KNeighborsClassifier(n_neighbors=9, p=2,metric='minkowski')
```

```
model1.fit(x_os,y_os)
```

```
KNeighborsClassifier(n_neighbors=9)
```

```
model1.score(x_os,y_os)
```

```
0.8132992327365729
```

```
y_pred1=model1.predict(x_test)
```



```

model.score(x_test,y_test)

0.7718120805369127

from sklearn.metrics import confusion_matrix,classification_report
cm=confusion_matrix(y_test,y_pred1)
cm

array([[77, 22],
       [12, 38]], dtype=int64)

plt.figure(figsize=(6,4))
sns.heatmap(confusion_matrix(y_test,y_pred1),annot=True,fmt='0.0f');

```



```

print(classification_report(y_test, y_pred1))

```

	precision	recall	f1-score	support
0	0.87	0.78	0.82	99
1	0.63	0.76	0.69	50
accuracy			0.77	149
macro avg	0.75	0.77	0.76	149
weighted avg	0.79	0.77	0.78	149

```

sensitivity=cm[1,1]/(cm[1,1]+cm[1,0])
specificity=cm[0,0]/(cm[0,0]+cm[0,1])
print("Sensitivity:", sensitivity)
print("Specificity:", np.round((specificity),2))

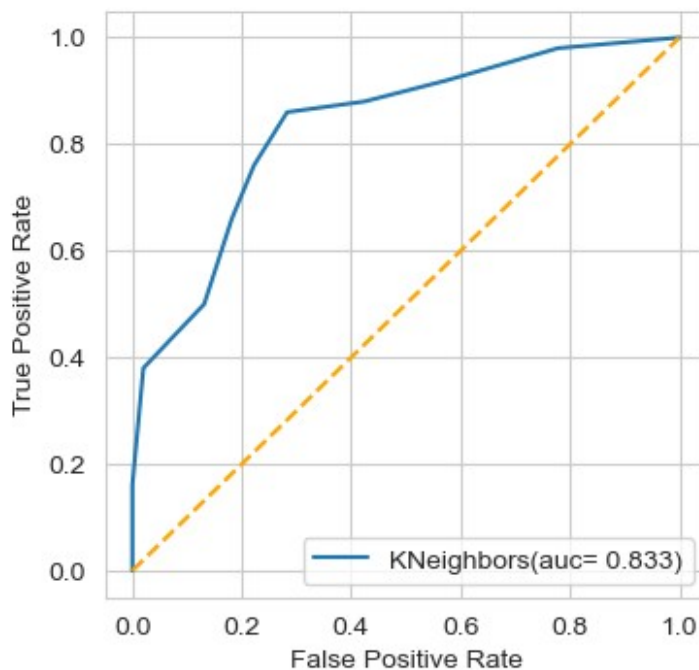
```

Sensitivity: 0.76
Specificity: 0.78

ROC Curve & AUC

```
from sklearn.metrics import roc_curve, auc
y_curve=model1.predict_proba(x_test)[:,-1]
FPR,TPR,threshold=roc_curve(y_test,y_curve)
auc=auc(FPR,TPR)

plt.figure(figsize=(4,4),dpi=100)
plt.plot(FPR,TPR,label="KNeighbors(auc= %0.3f)"% auc)
plt.plot([0,1],[0,1],color='orange',linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



Model Building Using Random Forest Classifier

```
model3=RandomForestClassifier(n_estimators=70,max_depth=20,criterion='entropy',min_samples_split=4,
```

```
max_leaf_nodes=25,min_samples_leaf=3,max_samples=0.2,max_features=3)
```

```
model3.fit(x_os,y_os)
```

```
RandomForestClassifier(criterion='entropy', max_depth=20,  
max_features=3,
```

```
max_leaf_nodes=25, max_samples=0.2,
```

```
min_samples_leaf=3,
```

```
min_samples_split=4, n_estimators=70)
```

```

model3.score(x_test,y_test)

0.7651006711409396

model3.score(x_os,y_os)

0.8337595907928389

y_pred3=model3.predict(x_test)

```

Classification metrics

```

from sklearn.metrics import confusion_matrix, classification_report
confusion_matrix(y_test,y_pred3)
cm=confusion_matrix(y_test,y_pred3)

plt.figure(figsize=(6,4))
sns.heatmap(confusion_matrix(y_test, y_pred3), annot = True, fmt =
'0.0f')

```

<AxesSubplot:>



```

print(classification_report(y_test, y_pred3))

```

	precision	recall	f1-score	support
0	0.89	0.74	0.81	99
1	0.61	0.82	0.70	50
accuracy			0.77	149
macro avg	0.75	0.78	0.75	149
weighted avg	0.80	0.77	0.77	149

```
sensitivity=cm[1,1]/(cm[1,1]+cm[1,0])
specificity=cm[0,0]/(cm[0,0]+cm[0,1])
print("Sensitivity:", sensitivity)
print("Specificity:", np.round((specificity),2))
```

Sensitivity: 0.82

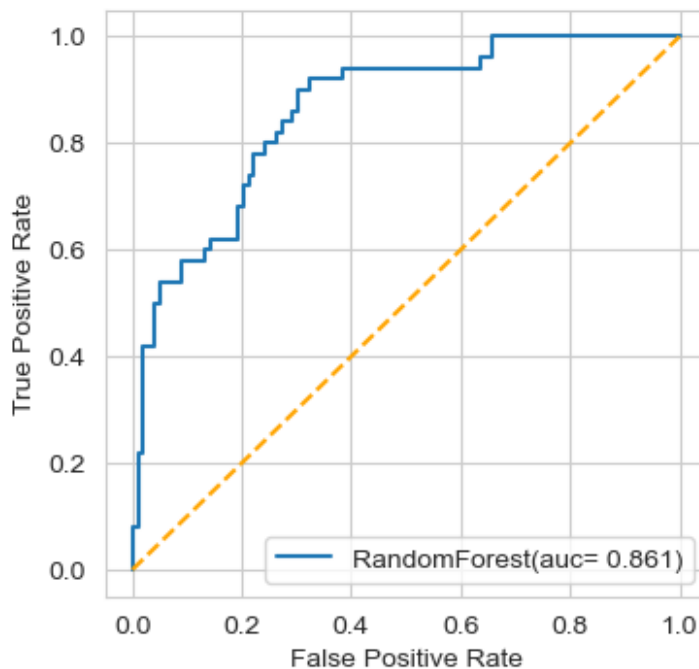
Specificity: 0.74

Roc curve & AUC

```
from sklearn.metrics import roc_curve,auc
y_curve=model3.predict_proba(x_test)[:,-1]
```

```
FPR,TPR,threshold=roc_curve(y_test,y_curve)
auc_RF=auc(FPR,TPR)
```

```
plt.figure(figsize=(4,4),dpi=100)
plt.plot(FPR,TPR,label="RandomForest(auc= %0.3f)"% auc_RF)
plt.plot([0,1],[0,1],color='orange',linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



Model Building Using Logistic Regression

```
from sklearn.linear_model import LogisticRegression
model2=LogisticRegression(solver = 'newton-cg', C = 1000,
penalty='l2')
```

```
model2.fit(x_os,y_os)
```

```
LogisticRegression(C=1000, solver='newton-cg')
```

```
model2.score(x_os,y_os)
```

```
0.7506393861892583
```

```
model2.score(x_test,y_test)
```

```
0.7583892617449665
```

```
y_pred2=model2.predict(x_test)
```

Classification Metrics

```
from sklearn.metrics import confusion_matrix, classification_report  
confusion_matrix(y_test,y_pred2)  
cm=confusion_matrix(y_test,y_pred2)
```

```
plt.figure(figsize=(6,4))
```

```
sns.heatmap(confusion_matrix(y_test, y_pred2), annot = True, fmt =  
'0.0f')
```

<AxesSubplot:>



```
print(classification_report(y_test, y_pred2))
```

	precision	recall	f1-score	support
0	0.87	0.75	0.80	99
1	0.61	0.78	0.68	50
accuracy			0.76	149
macro avg	0.74	0.76	0.74	149

weighted avg 0.78 0.76 0.76 149

```
sensitivity=cm[1,1]/(cm[1,1]+cm[1,0])
specificity=cm[0,0]/(cm[0,0]+cm[0,1])
print("Sensitivity:", np.round((sensitivity),2))
print("Specificity:", np.round((specificity),2))
```

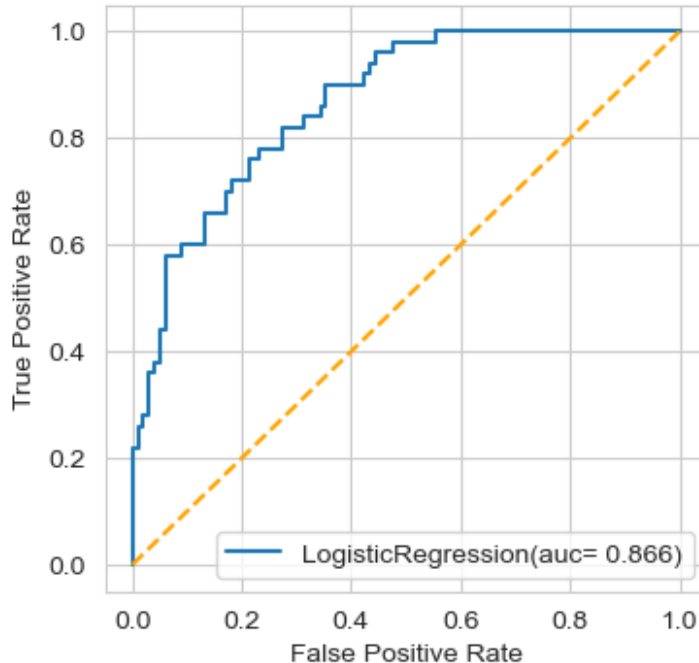
Sensitivity: 0.78
Specificity: 0.75

ROC Curve & AUC

```
from sklearn.metrics import roc_curve,auc
y_curve=model2.predict_proba(x_test)[:,-1]
```

```
FPR,TPR,threshold=roc_curve(y_test,y_curve)
auc=auc(FPR,TPR)
```

```
plt.figure(figsize=(4,4),dpi=100)
plt.plot(FPR,TPR,label="LogisticRegression(auc= %0.3f)"% auc)
plt.plot([0,1],[0,1],color='orange',linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



Model Building Using SVM

```
from sklearn.svm import SVC
model4=SVC(C=10,gamma=0.0001,probability=True)
```

```

model4.fit(x_os,y_os)
SVC(C=10, gamma=0.0001, probability=True)
model4.score(x_os,y_os)
0.7800511508951407
model4.score(x_test,y_test)
0.7449664429530202
y_pred4=model4.predict(x_test)

```

Classification Metrics

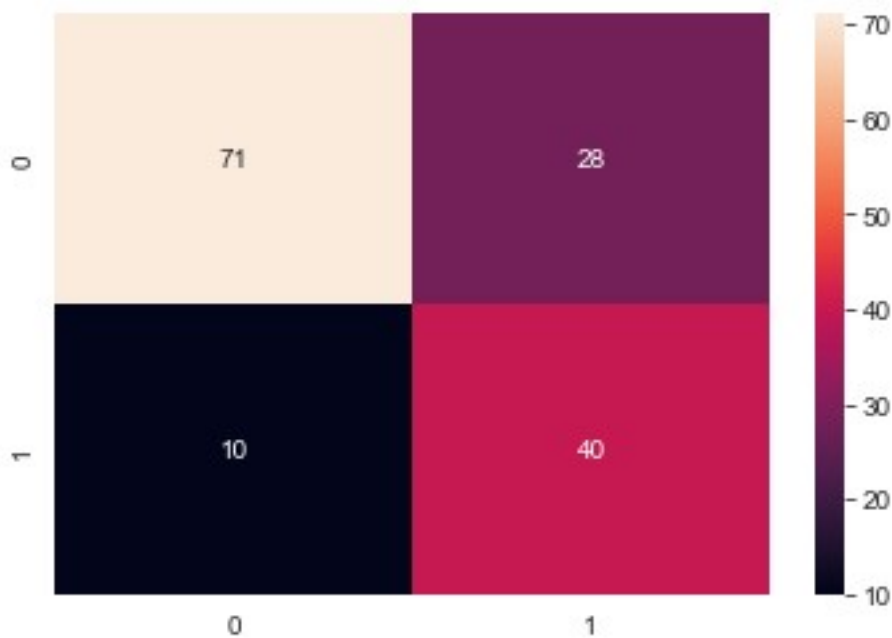
```

from sklearn.metrics import confusion_matrix, classification_report
confusion_matrix(y_test,y_pred4)
cm=confusion_matrix(y_test,y_pred4)

plt.figure(figsize=(6,4))
sns.heatmap(confusion_matrix(y_test, y_pred4), annot = True, fmt =
'0.0f')

```

<AxesSubplot:>



```

print(classification_report(y_test, y_pred4))

```

	precision	recall	f1-score	support
0	0.88	0.72	0.79	99
1	0.59	0.80	0.68	50

accuracy			0.74	149
macro avg	0.73	0.76	0.73	149
weighted avg	0.78	0.74	0.75	149

```
sensitivity=cm[1,1]/(cm[1,1]+cm[1,0])
specificity=cm[0,0]/(cm[0,0]+cm[0,1])
print("Sensitivity:", np.round((sensitivity),2))
print("Specificity:", np.round((specificity),2))
```

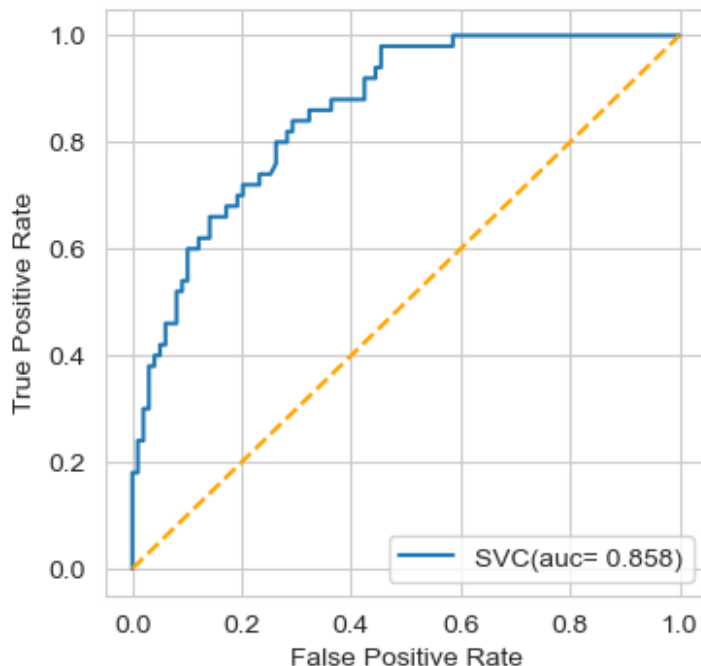
Sensitivity: 0.8
Specificity: 0.72

Roc Curve & AUC

```
from sklearn.metrics import roc_curve, auc
y_curve=model4.predict_proba(x_test)[: ,1]
```

```
FPR,TPR,threshold=roc_curve(y_test,y_curve)
auc=auc(FPR,TPR)
```

```
plt.figure(figsize=(4,4),dpi=100)
plt.plot(FPR,TPR,label="SVC(auc= %0.3f)"% auc)
plt.plot([0,1],[0,1],color='orange',linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



Model Building Using Naive Bayes

```
from sklearn.naive_bayes import GaussianNB  
model5=GaussianNB(var_smoothing=1e-07)
```

```
model5.fit(x_os,y_os)
```

```
GaussianNB(var_smoothing=1e-07)
```

```
model5.score(x_os,y_os)
```

```
0.7468030690537084
```

```
model5.score(x_test,y_test)
```

```
0.7516778523489933
```

```
y_pred5=model5.predict(x_test)
```

Classification Metrics

```
from sklearn.metrics import confusion_matrix, classification_report  
confusion_matrix(y_test,y_pred5)  
cm=confusion_matrix(y_test,y_pred5)
```

```
plt.figure(figsize=(6,4))
```

```
sns.heatmap(confusion_matrix(y_test, y_pred5), annot = True, fmt =  
'0.0f')
```

<AxesSubplot:>



```
print(classification_report(y_test, y_pred5))
```

	precision	recall	f1-score	support
0	0.87	0.74	0.80	99
1	0.60	0.78	0.68	50
accuracy			0.75	149
macro avg	0.73	0.76	0.74	149
weighted avg	0.78	0.75	0.76	149

```
sensitivity=cm[1,1]/(cm[1,1]+cm[1,0])
specificity=cm[0,0]/(cm[0,0]+cm[0,1])
print("Sensitivity:", np.round((sensitivity),2))
print("Specificity:", np.round((specificity),2))
```

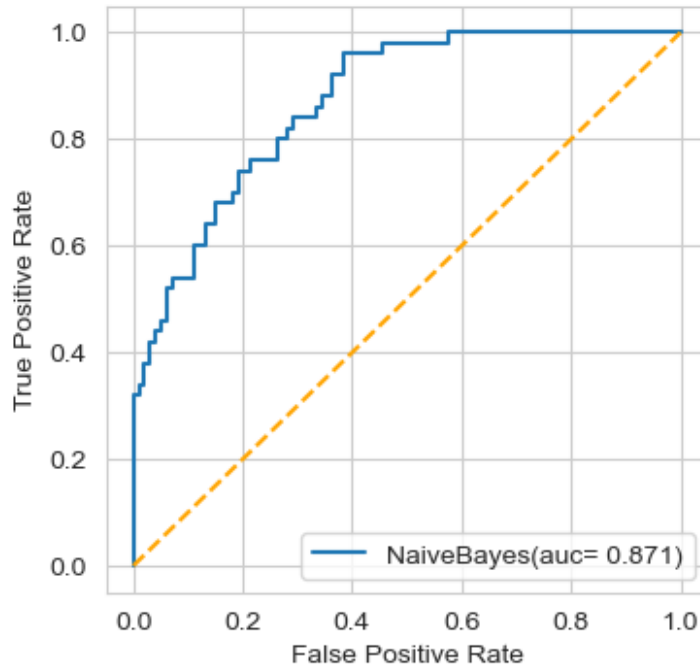
Sensitivity: 0.78
Specificity: 0.74

ROC Curve & AUC

```
from sklearn.metrics import roc_curve,auc
y_curve=model5.predict_proba(x_test)[:,-1]
```

```
FPR,TPR,threshold=roc_curve(y_test,y_curve)
auc=auc(FPR,TPR)
```

```
plt.figure(figsize=(4,4),dpi=100)
plt.plot(FPR,TPR,label="NaiveBayes(auc= %0.3f)"% auc)
plt.plot([0,1],[0,1],color='orange',linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```




Comparing the results of all the models taking into account F1-score,precision,accuracy & AUC.

Considering AUC , Naive Baiyes Tops the list.

considering F1-score,precision,accuracy, Random Forest Classifier is at the top.

Here, I choose Random Forest Classifier as the best model.

Pie Chart



Count of Outcome


768

Outcome Label

Diabetic

Non-Diabetic


Pregnancy vs BP




BP vs Glucose




SkinThickness & BMI




Age & DPF




Histogram BMI




Histogram Insulin




Histogram Glucose




Histogram BP




Heatmap BMI & SkinThickness




Bubble showing Glucose vs age




Pregnancy vs BP



SkinThickness & BMI




BP vs Glucose




Age & DPF




Heatmap BMI & SkinThickness




Histogram Insulin




Histogram BP



Histogram Glucose



Histogram BMI



Bubble showing Glucose vs age



Pie Chart

